

Seamless Switch Between Local and Cloud APIs

Source: [Ollama, 2025] [OpenAI, 2024]

Local Deployment

Zero Cost · Data Private · No Rate Limit

```
base_url = "http://localhost:11434/v1"
```

```
api_key = "ollama" · model = "qwen3.5:2b"
```

Pros

- ❑ Free experimentation
- ❑ Data stays local
- ❑ No rate limits
- ❑ Full control
- ❑ Limited capability
- ❑ Speed depends on hardware

Cons

- ❑ Requires GPU
- ❑ Limited concurrency
- ❑ Limited model selection
- ❑ No SLA
- ❑ High maintenance
- ❑ Not for production

Just
2 lines

Cloud API

Pay-per-use · Best Performance · Rate Limited

```
base_url = "https://api.openai.com/v1"
```

```
api_key = "sk-..." · model = "gpt-4o"
```

Pros

- ❑ Best performance
- ❑ No GPU needed
- ❑ High concurrency SLA
- ❑ More model choices
- ❑ Pay per token
- ❑ Data sent to cloud

Cons

- ❑ Cost scales with usage
- ❑ Rate limits
- ❑ Parameter limits
- ❑ Network dependent
- ❑ Not customizable
- ❑ Privacy risk