

Positional Encoding Comparison

Reference: [Su et al., 2024] [Press et al., 2022] [Chen et al., 2023]

Absolute Positional Encoding

$$h = x + \text{PE}(\text{pos})$$

Sinusoidal / Learnable vectors
Added directly to embeddings

- ❑ Simple implementation
- ❑ Cannot extrapolate to longer sequences
- ❑ Performance drops beyond training length

Used by: Original Transformer

BERT, T5

Extrapolation: Poor

Current Status: Deprecated

RoPE (Rotary Positional Encoding)

Rotate Query and Key vectors

$$\text{rot} = f(|\text{pos}_i - \text{pos}_j|)$$

Depends only on relative position difference

- ❑ Naturally supports relative position
- ❑ Good extrapolation ability
- ❑ Linear attention compatible

Used by: LLaMA 2/3

Qwen, Mistral

Extrapolation: Good

Current Status: Mainstream

ALiBi (Linear Bias)

$$\text{bias} += m \times (j - i)$$

Does not modify embeddings
Added directly to attention matrix

- ❑ Train 2K, infer 64K
- ❑ Best extrapolation ability
- ❑ Requires tuning slope m

Used by: MPT

BLOOM

Extrapolation: Best

Current Status: Niche