


# RadixAttention Radix Prefix Cache Tree

Data source: [Zheng et al., 2024]



RadixAttention uses a prefix tree to organize KV Cache; requests with shared prefixes reuse cached data, reducing KV Cache memory by 67-85%